

CS 388P Writeup

Jesus Gonzalez, Tomas McCandless

October 9, 2012

jsdklshasdsf [?] [?]

Models of parallel computation should attempt to satisfy two goals that are in conflict: a model should be sufficiently abstract to allow algorithm designers to write programs that are simple and portable across architectures, and a model should also expose some low-level architectural details to allow for optimization. The PRAM model is both widely used and simple, yet it has been criticized for being too high-level and thus failing to accurately model parallel machines. Specifically, the PRAM does not model realities of current parallel machines, such as bandwidth limitations. Similarly, network-based models such as the hypercube have been criticized for being too low-level, failing to be widely reflect the current technological state of parallel machines.

0.1 Emulations between QSM, BSP and LogP: A framework for general-purpose parallel algorithm design.

[?]

The authors present work-preserving emulations with just a small slowdown between the geneal-purpose models: LogP, BSP and QSM. The relevance of

this lies on the simplicity of the QSM model, which makes it easy to program for and, yet, through emulation on more realistic models like BSP and LogP run an algorithm with just a logarithmic slowdown.

The authors present 3 algorithms for basic problems (prefix sums, sample sort and list ranking). All of them are optimal and simulations show that the predicted performance for each of them is accurate.

0.2 Emulations between QSM, BSP and LogP: A framework for general-purpose parallel algorithm design.

[?]

PRAM is a very simple model for parallel computer that abstracts too much parameters from real machines, which turns into apparently efficient models that turn to run slowly on real machines. Therefore, exists the need of more realistic models that reflect with more fidelity a machine and yet is not too complicated so that implementation is kept simple.

0.2.1 Brief description of the models

All the next models have work measured as the time-processors product.

BSP BSP (Bulk-Synchronous Parallel) takes 3 parameters:

- p processor/memory components: Components that communicate via point-to-point messages.
- Bandwidth g .

- Latency L .

A computation in this model consists of a sequence of supersteps, each separated by a bulk synchronization. In each superstep, a component can perform local computations, send and receive a set of messages. The cost T superstep is calculated as $T = \max(w, gh, L)$ and the overall time that takes an algorithm to run is the sum of T of every superstep.

LogP LogP (Latency, overhead, gap, Processors) takes 4 parameters and a derived value:

- Processors p : Number of processors.
- Latency l : Time taken by network to transmit a message from one processor to another is at most l .
- Gap g : Least units of time taken between sending or reception of a message by a processor.
- Overhead o : Time taken by a processor to send or receive a message.
- Capacity constraint: Maximum number of messagings heading to the same processor can't be larger than $\lceil l/g \rceil$

Once a processor has reached its capacity, another processor trying to send a message will stall. The time taken by a LogP algorithm is the longest time taken by any processor to end its operation.

QSM QSM (Queuing Shared Memory) consists of p processors that communicate through shared memory, but run computations on private memory.

An algorithm consists of the analogous for a superstep on BSP, a synchronized phase. Concurrent reads and concurrent writes are permitted on a single synchronized phase over the same shared-memory location, but not both.

Metrics:

- Processors p .
- Maximum contention k : of a synchronized phase Largest number of processors reading or writing a location.

The time cost of a phase is $\max(m_{op}, gm_{rw}, k)$, where m_{op} = maximum operations performed by any processor and m_{rw} = maximum read/write operations performed by any processor. The time taken by an algorithm is the sum of the cost of all of its phases.

s-QSM is a variation of QSM where the time cost of a phase is $\max(m_{op}, gm_{rw}, gk)$

aoeuaou

References

- [1] P. Gibbons, Y. Matias, V. Ramachandran. "Can a shared memory model serve as a bridging model for parallel computation?" Theory of Computing Systems Special Issue on papers from SPAA'97, vol. 32, no. 3, 1999, pp. 327-359
- [2] V. Ramachandran, B.Grayson, M. Dahlin.(2003) Emulations between QSM, BSP and LogP: A framework for general-purpose parallel algorithm design. Journal of Parallel and Distributed Computing, vol. 63, 2003, pp. 1175-1192.